

# SVM Dual Problem and Kernel Methods

Yaniv Tal

18-Jun-2021



# Agenda

- SVM – The Primal problem
- SVM – The Dual problem
- Linear Inseparability and Kernels
- Common Kernels
- Implementation
- Results on Datasets

The background is a complex digital-themed abstract. It features a dark blue base with lighter blue lines and dots forming a network or neural structure. Binary code (0s and 1s) is scattered throughout. A faint, light blue silhouette of a human head is visible on the right side. A solid orange horizontal bar is positioned above the text.

# SVM Primal Problem

# SVM – The Primal Problem

- Linear separator, optimizes for largest margin

We have  $n$  data points,  $p$  features


$$x_i, w \in \mathbb{R}^p$$

$$b \in \mathbb{R}$$

$$y_i \in \{-1, 1\}$$

$$x \in \mathbb{R}^{n \times p}$$

$$y \in \mathbb{R}^n$$



# SVM – The Primal Problem


- Margin size formulation :

Let  $x_0$  be a point on the hyperplane, so:  $w^T x_0 - b = 0$ .

If we move from  $x_0$  a distance  $k$  in the direction of  $w$  to the hyperplane  $w^T x - b = 1$ ,  
we get a point:  $x_0 + k \cdot \frac{w}{\|w\|}$ .

The new point is on the upper plane, so:  $w^T \cdot (x_0 + k \cdot \frac{w}{\|w\|}) - b = 1$

And so we get:  $k = \frac{1}{\|w\|}$



# SVM – The Primal Problem

- Margin size formulation :

Let  $x_0$  be a point on the hyperplane, so:  $w^T x_0 - b = 0$ .


If we move from  $x_0$  a distance  $k$  in the direction of  $w$  to the hyperplane  $w^T x - b = 1$ , we get a point:  $x_0 + k \cdot \frac{w}{\|w\|}$ .

The new point is on the upper plane, so:  $w^T \cdot (x_0 + k \cdot \frac{w}{\|w\|}) - b = 1$

And so we get:  $k = \frac{1}{\|w\|}$

► Hence, the margin we want to maximize is:  $\frac{2}{\|w\|}$

► Which is the same as minimizing  $\frac{1}{2} \|w\|^2$



# SVM – The Primal Problem

- **Constraints:** All points should be on the correct sides of the lines:


$$y_i(x_i^T w + b) \geq 1$$
$$i = 1..n$$

- And so, we get the primal SVM problem:

$$\min_{w,b} \frac{1}{2} \|w\|^2$$

**s.t.**

$$y_i(x_i^T w + b) \geq 1$$
$$i = 1..n$$



# SVM – The Primal Problem

- Primal SVM problem:

$$\min_{w,b} \frac{1}{2} \|w\|^2$$

**s.t.**


$$y_i(x_i^T w + b) \geq 1$$

$$i = 1..n$$

- Lagrangian:

$$\mathcal{L}_p(w, b, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^n \alpha_i (y_i(x_i^T w + b) - 1)$$

- Quadratic problem, solved with KKT conditions
- Note that the only active constraints are the points that define the margins -  
from KKT – all other Alphas are zero!





The background is a blue-toned abstract image. It features a faint silhouette of a human head in profile, facing right. Overlaid on this are various digital elements: binary code (0s and 1s) scattered throughout, a network of glowing blue lines and nodes resembling a neural network or data flow, and some faint, larger-scale grid patterns. A solid orange horizontal bar is positioned across the middle of the image, just above the text.

# SVM Dual Problem

# SVM – The Dual Problem (Strong duality)

## Formulation:

$$\text{Max}_{\alpha} \text{Min}_{w,b} \mathcal{L}(w, b, \alpha)$$

$$\mathcal{L}(w, b, \alpha) = \frac{1}{2} w^T w - \sum_{i=1}^n \alpha_i (y_i (w^T x_i - b) - 1)$$

$$g(\alpha) = \text{Min}_{w,b} \mathcal{L}(w, b, \alpha)$$

$$\frac{\partial \mathcal{L}}{\partial w} = 0 \implies w = \sum_i \alpha_i y_i x_i$$


$$\frac{\partial \mathcal{L}}{\partial b} = 0 \implies \alpha^T y = 0$$

$$g(\alpha) = \sum_{i=0}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j, \quad \alpha^T y = 0, \quad \alpha_i \geq 0$$

$$\text{Max}_{\alpha} g(\alpha)$$

**s.t.**

$$\alpha^T y = 0, \quad \alpha_i \geq 0$$



# SVM – The Dual Problem (Strong duality)

## Formulation:

$$\text{Max}_{\alpha} g(\alpha)$$

**s.t.**

$$\alpha^T y = 0, \quad \alpha_i \geq 0$$

► Which is same as:

$$\text{Min}_{\alpha} \left( \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j x_i^T x_j \right) - \sum_{i=1}^n \alpha_i$$

**s.t.**


$$\alpha^T y = 0, \quad \alpha_i \geq 0$$

► Remember - Alpha=0 for non-support vectors

$$\text{Min}_{\alpha} \left( \frac{1}{2} \sum_{SV \text{ pairs}(i,j)} \alpha_i \alpha_j y_i y_j x_i^T x_j \right) - \sum_{SV \text{ alphas}} \alpha_i$$


**s.t.**


$$\alpha^T y = 0, \quad \alpha_i \geq 0$$



# Why go Dual?


- “Wide” datasets - High dimensional data ( $p > n$ ):
  - ▶ Primal problem size –  $n \times p$  (Size of  $X$ )
  - ▶ Dual problem size –  $n \times n$  (Size of  $X^T X$ )
- Kernels!



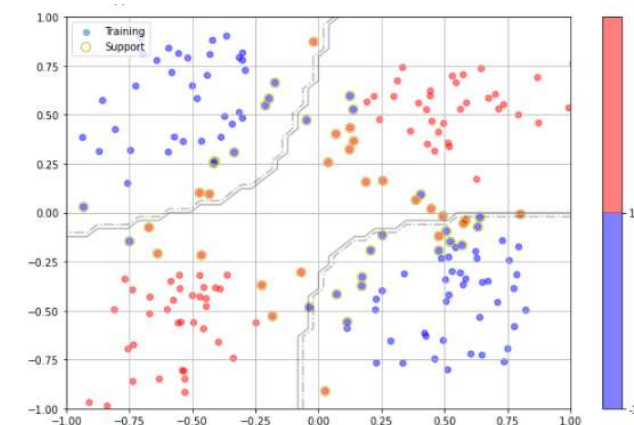

The background is a dark blue field filled with glowing white and light blue elements. On the left, there are faint, semi-transparent outlines of human heads facing right. Overlaid on these and the background are various patterns: a grid of binary digits (0s and 1s), a network of interconnected nodes and lines resembling a neural network or data flow, and some larger, fainter binary strings. A solid orange horizontal bar is positioned above the text area.

# SVM with Nonlinearly Separable Data, Kernel Functions

# Non-linearly separable data



- Cover's theorem – Linear separability is more likely in higher dimensions
- The good news: Given enough dimensions, Everything is linearly separable!
- The bad news: We get a VERY wide matrices....



# Adding Features the Standard Way

- Reminder – The dual problem:

$$\begin{aligned} & \text{Min}_{\alpha} \left( \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \underbrace{x_i^T x_j} \right) - \sum_{i=1}^n \alpha_i \\ & \text{s.t.} \\ & \alpha^T y = 0, \quad \alpha_i \geq 0 \end{aligned}$$

- Define:

$$\begin{aligned} & \varphi(x) : \mathbb{R}^p \rightarrow \mathbb{R}^P, \quad P \geq p \\ & k(x, y) = \varphi(x) \cdot \varphi(y), \quad k(x, y) : \mathbb{R}^p \times \mathbb{R}^p \rightarrow \mathbb{R} \end{aligned}$$

- So, the problem becomes:

$$\begin{aligned} & \text{Min}_{\alpha} \left( \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \underbrace{k(x_i, x_j)} \right) - \sum_{i=1}^n \alpha_i \\ & \text{s.t.} \\ & \alpha^T y = 0, \quad \alpha_i \geq 0 \end{aligned}$$

- Dot product between vectors of dimension P instead of p.



# Adding features – the Kernel Trick

- A function  $k : \mathbb{R}^p \times \mathbb{R}^p \rightarrow \mathbb{R}$  is called a kernel if there exists a mapping function  $\varphi : \mathbb{R}^p \rightarrow \mathbb{R}^P$  so that the following always holds:  $\forall x, y : k(x, y) = \varphi(x) \cdot \varphi(y)$
- A mapping function  $\varphi : \mathbb{R}^p \rightarrow \mathbb{R}^P$  is said to afford a kernel if such  $k$  exists
- Example – Inhomogeneous polynomial kernel of degree 2:

$$x, y \in \mathbb{R}^2$$

$$k(x, y) = (x^T y + 1)^2$$

$$\varphi(x) = [x_1^2, \sqrt{2}x_1, \sqrt{2}x_1x_2, \sqrt{2}x_2, x_2^2, 1]$$

$$\varphi : \mathbb{R}^2 \rightarrow \mathbb{R}^6$$

$$\varphi(x) \cdot \varphi(y) : \mathbb{R}^6 \times \mathbb{R}^6 \rightarrow \mathbb{R}$$

$$k : \mathbb{R}^2 \times \mathbb{R}^2 \rightarrow \mathbb{R}$$

- Less computation, same result!





# Some Common Useful Kernels

- Homogeneous Polynomial kernel:  $k(x, y) = (x^T y)^d$
- Inhomogeneous Polynomial kernel:  $k(x, y) = (x^T y + c)^d$
- Radial Basis Function (RBF) kernel:  $k(x, y) = e^{-\frac{\|x - y\|^2}{2\sigma^2}}$





Fig 6: RBF Kernel SVM for Iris Dataset [Image Credits: <https://scikit-learn.org/>]



## Additional notes

- Some kernel arithmetic:

$$(1) \ k(x, y) = k_1(x, y) + k_2(x, y)$$

$$(2) \ k(x, y) = ak_1(x, y) \text{ where } a > 0$$

$$(3) \ k(x, y) = f(x) \cdot f(y) \text{ for any function } f \text{ on } x$$

$$(4) \ k(x, y) = k_1(x, y) \cdot k_2(x, y)$$

$$(5) \ k(x, y) = \frac{k_1(x, y)}{\sqrt{k_1(x, x)}\sqrt{k_1(y, y)}}$$

- Not all kernels are useful

The background is a deep blue with a complex, abstract pattern. It features glowing white and light blue lines that form a network or circuit-like structure. Interspersed throughout are strings of binary code (0s and 1s) in a lighter blue font. The overall effect is a high-tech, digital aesthetic.

# Implementation

# Implementation – SVM Kernel Classifier

- Classifier class usage example:

```
print(f'[Main] Info: Building SVM model with C={C}, sigma={sigma}')
model = KernelSvmClassifier(C=C, kernel=RBF)
model.fit(X_train, y_train, DEBUG=DEBUG)

print('[Main] Info: Test model on test data')
test_prediction = model.predict(X_test, DEBUG=DEBUG)

print('[Main] Info: Finished prediction! Test results:')
run_results_df = data_helpers.assess_accuracy(test_prediction, y_test)
```

- Kernel function:

```
def RBF(x1, x2):
    """RBF kernel"""
    diff = x1 - x2
    return np.exp(-1 * (diff @ diff) / (2 * (sigma**2)))
```

```
def poly(x1, x2):
    """Poly kernel"""
    return (x1 @ x2 + 1) ** poly_rank
```

# Implementation – SVM Kernel Classifier (Constrained)

- Fit() method – cost function:

```
# Create Gram matrix of k(x) y:
gramXX = np.apply_along_axis(lambda x1: np.apply_along_axis(lambda x2: self.kernel(x1, x2), 1, X), 1, X) # 2 for loops.
yp = y.reshape(-1, 1)
yy = yp @ yp.T
gramXXyy = gramXX * yy
```

```
# Lagrange dual objective function (to maximize!)
def ld_obj(gram, alpha):
    return alpha.sum() - 0.5 * alpha @ (alpha @ gram)

# Partial derivative of ld_obj on alpha
def d_ldobj_d_alpha(gram, alpha):
    return np.ones_like(alpha) - alpha @ gram

# cost function and its gradient - to minimizes
def cost_func(a):
    return -ld_obj(gramXXyy, a)

def cost_grad(a):
    return -d_ldobj_d_alpha(gramXXyy, a)
```

$$\text{Min}_{\alpha} \left( \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \underbrace{\alpha_i \alpha_j}_{\text{blue}} \underbrace{y_i y_j k(x_i, x_j)}_{\text{red}} \right) - \underbrace{\sum_{i=1}^n \alpha_i}_{\text{purple}}$$

Matrix writing:

$$\text{Min}_{\alpha} \left( \frac{1}{2} \underbrace{\alpha^T}_{\text{blue}} \underbrace{(\alpha^T \cdot \text{Gram}_k)}_{\text{red}} \right) - \underbrace{\alpha^T \cdot \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}}_{\text{purple}}$$

# Implementation – SVM Kernel Classifier (Constrained)

- Fit() method – constraints and (scipy) solver

```
# Constraints:
A = np.vstack((-np.eye(N), np.eye(N)))
b = np.hstack((np.zeros(N), self.C * np.ones(N)))
constraints = ({'type': 'eq', 'fun': lambda a: np.dot(a, y), 'jac': lambda a: y},      # y @ alpha = 0
               {'type': 'ineq', 'fun': lambda a: b - np.dot(A, a), 'jac': lambda a: -A}) # 0 <= alpha <= C
```

```
# Minimizing the negative dual
results = optimize.minimize(fun=cost_func,
                           x0=np.ones(N),
                           method='SLSQP',
                           jac=cost_grad,
                           constraints=constraints)
self.alpha = results.x
```

```
# store support vectors and y*alpha vals to use in inference
epsilon = 1e-8 # Max distance to become a support vector
supportIndices = self.alpha > epsilon
self.supportVectors = X[supportIndices]
self.supportAlphaY = y[supportIndices] * self.alpha[supportIndices]
```

## Implementation – SVM Kernel Classifier (Constrained)

- Predict() method – note it only uses the support vectors:

$$\text{Sign}\left(\sum_{SV} k(x, SV_i) \cdot \alpha_i y_i\right)$$

```
def predict(self, X, DEBUG=False):
    """ Predict y values in {-1, 1} """
    if DEBUG:
        print(f'[KernelSvmClassifier] Debug: predict() started -----')
        print(f'[KernelSvmClassifier] Debug: X.shape: {X.shape}')

    def predict_sample(x):
        x1 = np.apply_along_axis(lambda s: self.kernel(s, x), 1, self.supportVectors)
        # Calc kernel of the sample with each support vector, return vectorized results
        x2 = x1 * self.supportAlphaY # Multiply by alpha*y of each vector
        return np.sum(x2)

    d = np.apply_along_axis(predict_sample, 1, X) # for each vector in X, run predict_sample
    return 2 * (d > 0) - 1
```





Some Results



# Results on Actual Datasets – Epileptic Seizure Predictin

- Dataset:
  - Identify epileptic seizure based on brainwave data from 178 sensors
  - 178 attributes, binary target (Seizure / no seizure)
  - Train size: 800, Test size: 15,000
- Linear kernel with slack: **Failed to converge**

- RBF Kernel

Column1	C	sigma	% Success	% False neg	% False pos
3	0.1	0.1	89.64	0.28	10.08
10	1	0.1	89.64	0.28	10.08
17	10	0.1	89.64	0.28	10.08
18	10	0.5	85.08	0.06	14.86
11	1	0.5	82.79	0.03	17.19

- Polynomial kernel

Column1	C	Polynom rank	% Success	% False neg	% False pos
1	100	3	79.95	0.00	20.05
2	100	5	79.95	0.00	20.05
3	100	10	79.95	0.00	20.05
4	100	15	79.95	0.00	20.05
6	10	3	79.95	0.00	20.05

# Results on Actual Datasets – Network Attack Prediction

- Dataset:

- ▶ Predict whether a network transaction is an attack based on technical attributes – protocol, ports, packet attributes and more
- ▶ 41 attributes, binary target (attack true / false - **Union of all attack types**)
- ▶ Train size: 800, Test size: 15,000

- Linear kernel with slack:

Column1	C	Polynom rank	% Success	% False neg	% False pos
3	0.1		96.82	2.45	0.73
4	0.01		96.63	2.61	0.75
2	1		87.81	12.06	0.13
1	10		82.65	17.29	0.07
0	100		80.23	19.77	0.00

- RBF Kernel

Column1	C	sigma	% Success	% False neg	% False pos
24	1	0.1	99.21	0.42	0.37
31	10	0.1	99.21	0.43	0.35
18	0.1	0.5	99.18	0.04	0.78
26	1	0.75	99.17	0.11	0.71
34	10	1	99.05	0.14	0.81

- Polynomial kernel

Column1	C	Polynom rank	% Success	% False neg	% False pos
1	100	3	99.27	0.41	0.32
8	10	10	98.84	0.29	0.87
18	0.1	10	98.63	0.27	1.10
15	0.1	1	96.81	2.45	0.73
20	0.01	1	96.63	2.61	0.75

# Altered Digits Dataset – One vs All

## Linear kernel results:

# support vectors: (2896, 784)

Training set: {'Success': 1.0, 'False Positive': 0.0, 'False Negative': 0.0}

Test set: {'Success': 0.1864406779661017, 'False Positive': 0.8135593220338984, 'False Negative': 0.8135593220338984}

<Figure size 640x480 with 0 Axes>



[7 3 3 7 4 0 1 9 3 1 6 9 2 0 3 0 0 1 5 3 9 6 8 0 7]

## Polinomial kernel results:

# support vectors: (2054, 784)

Training set: {'Success': 0.8536666666666667, 'False Positive': 0.14633333333333334, 'False Negative': 0.14633333333333334}

Test set: {'Success': 0.7867381214902947, 'False Positive': 0.21326187850970538, 'False Negative': 0.21326187850970538}

<Figure size 640x480 with 0 Axes>



[4 9 1 1 8 9 1 4 8 1 6 1 0 4 4 2 5 9 4 1 4 8 0 8 7]

## rbf kernel results:

# support vectors: (2071, 784)

Training set: {'Success': 0.95, 'False Positive': 0.05, 'False Negative': 0.05}

Test set: {'Success': 0.9067411984922691, 'False Positive': 0.09325880150773097, 'False Negative': 0.09325880150773097}

<Figure size 640x480 with 0 Axes>



[4 9 1 7 8 9 1 4 8 1 6 2 0 4 4 2 5 9 4 2 4 8 4 8 7]